Supervision by token class

Speaker-diarized ASR training data

Share of supervised tokens

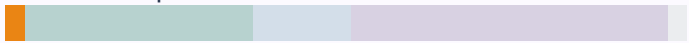

Share of unweighted loss

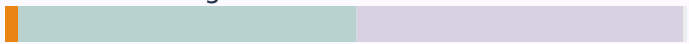

Token class	Count	Loss
Speaker ID	**2.90%**	**1.97%**
Timestamp	33.44%	49.51%
Formatting	14.48%	≈0.23%
Transcript	46.50%	47.74%
Other	2.68%	≈0.55%

Other: remaining supervised tokens

Reweight speaker-ID numbers, not formatting tokens

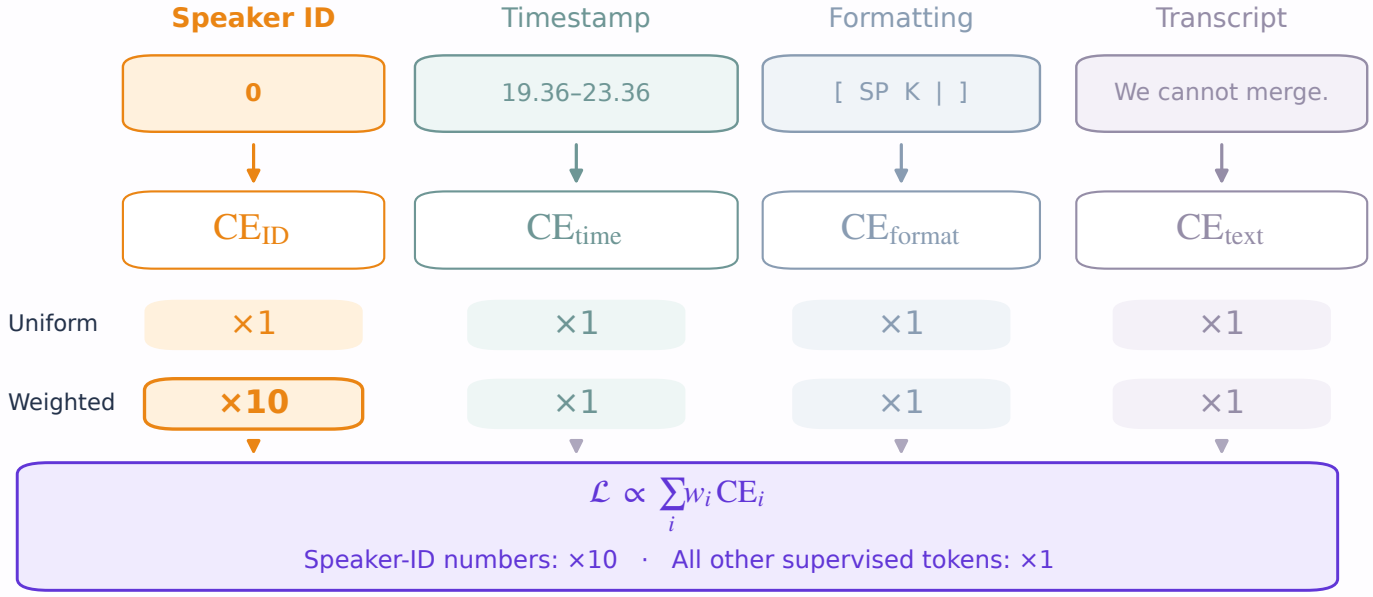

Target sequences and token counts remain unchanged.